

# Spark Structured Streaming

---

**Emanuele Della Valle**

---

prof @ Politecnico di Milano  
Co-founder @ Quantia Consulting

[emanuele.dellavalle@polimi.it](mailto:emanuele.dellavalle@polimi.it)

<http://emanueledellavalle.org>

# Overview

*Spark Structured Streaming provides fast, scalable, fault-tolerant, end-to-end exactly-once stream processing without the user having to reason about streaming.*

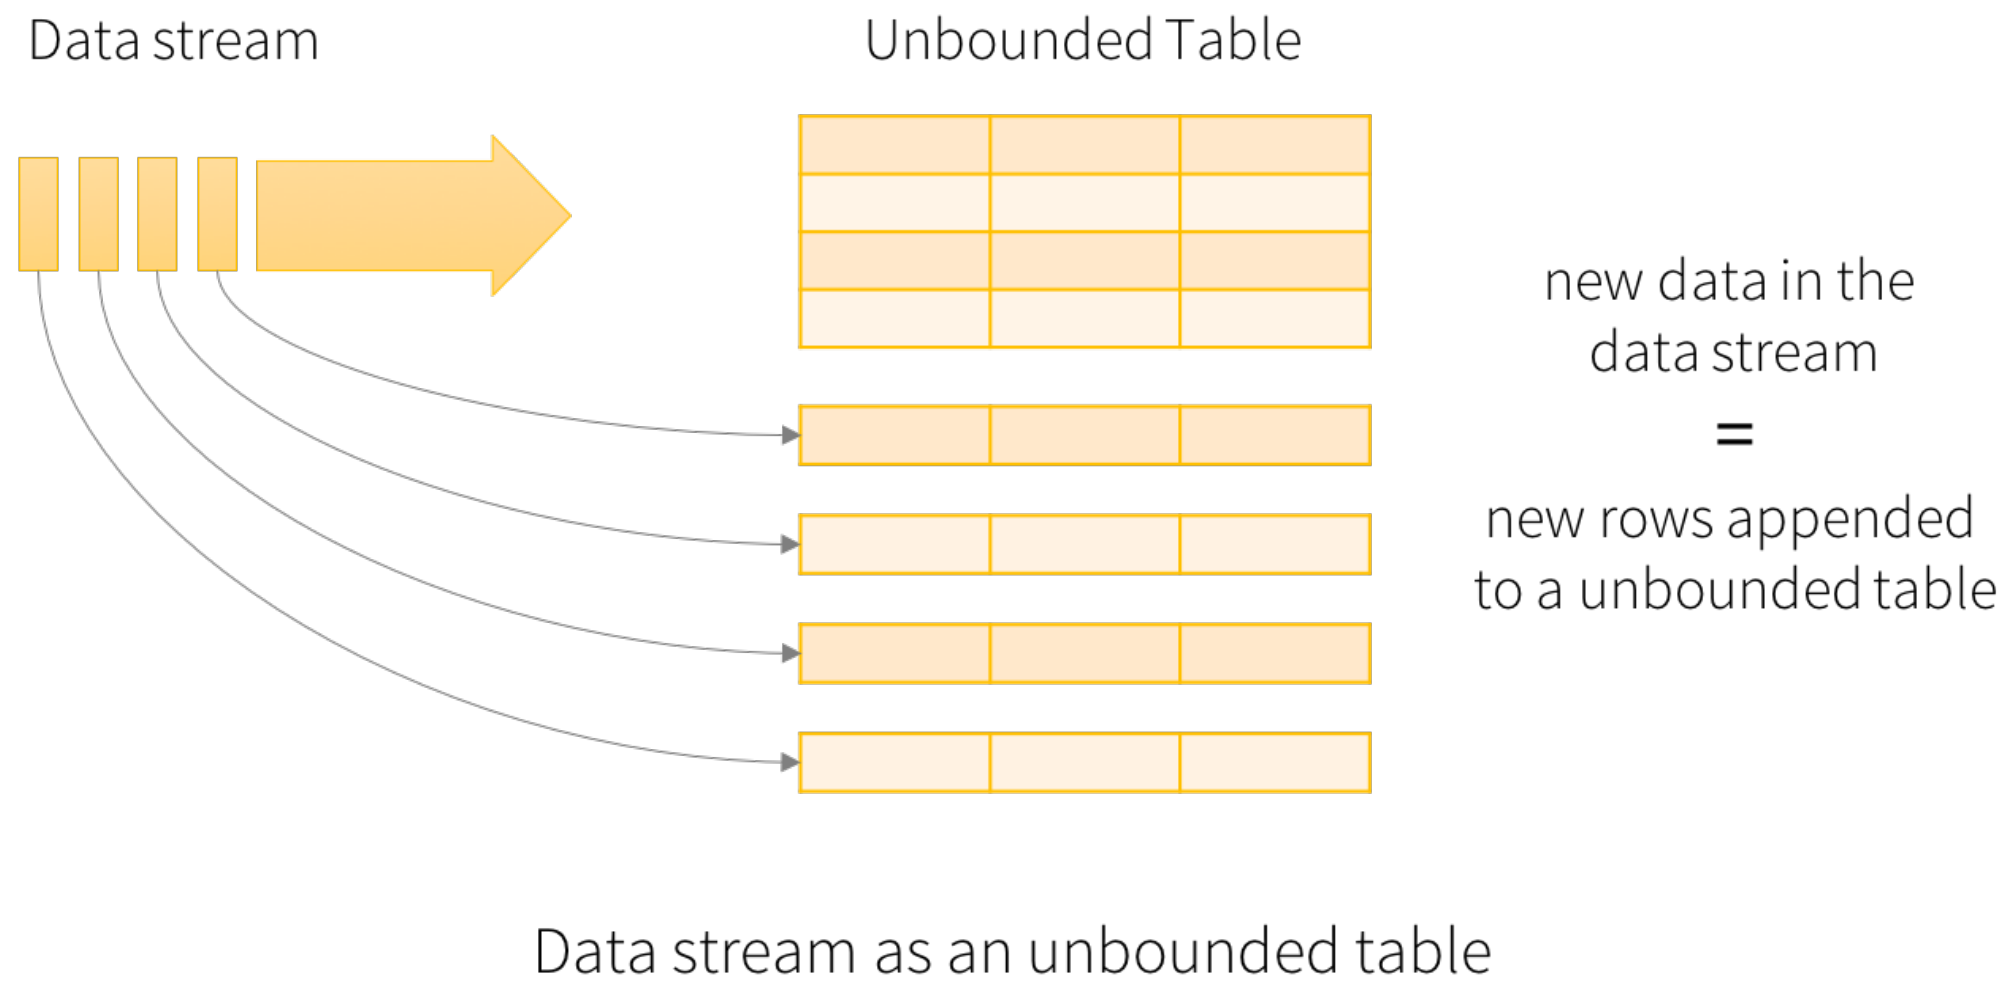
# Positioning Spark Structured Streaming

- default: **micro-batch processing** engine
  - end-to-end latencies as low as 100 milliseconds
  - exactly-once fault-tolerance guarantees
- since Spark 2.3, also **Continuous Processing**
  - end-to-end latencies as low as 1 millisecond
  - at-least-once guarantees
- choose the mode based on requirements

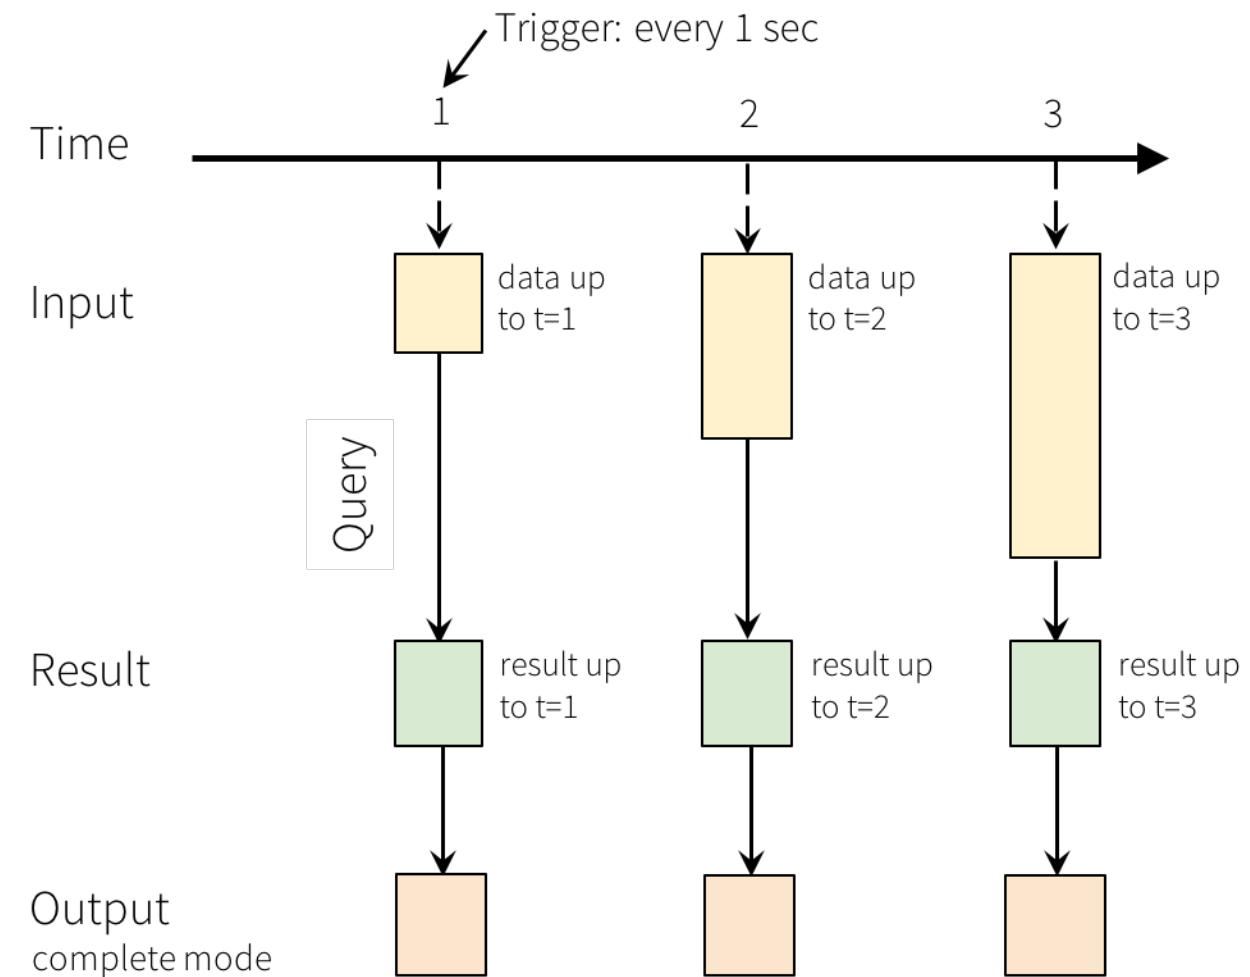
# Programming Model

- key idea: treat a live data stream as a table that is being continuously appended
- Result:
  - user can express streaming computation as standard batch-like query as on a static table
  - Spark runs it as an *incremental* query on the *unbounded* input table

# Basic Concepts 1/2



# Basic Concepts 2/2



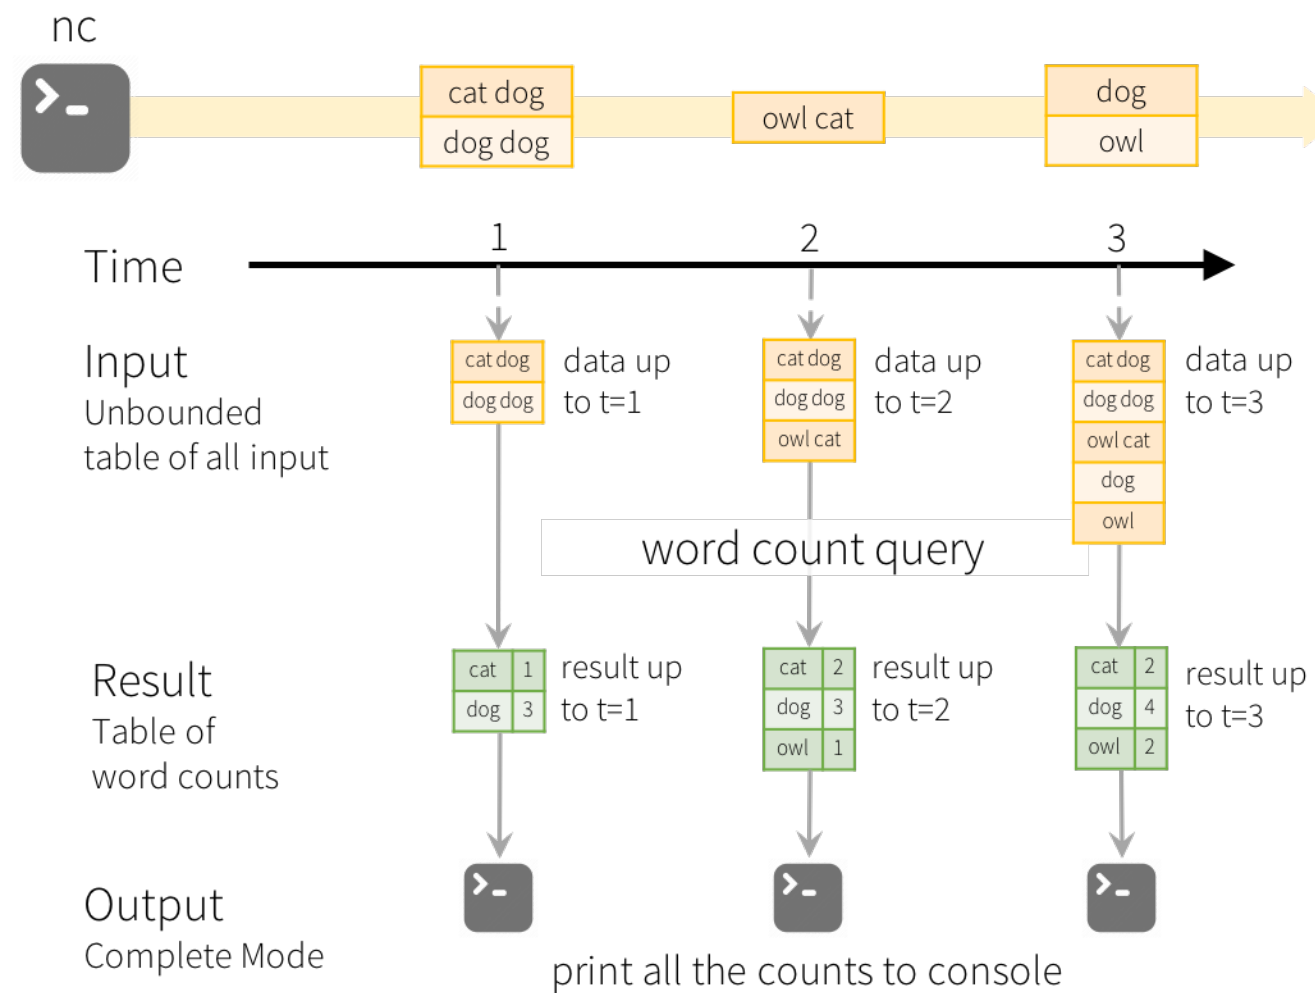
Programming Model for Structured Streaming

# Output Modes

- *Complete Mode*: The entire updated result is written
- *Append Mode*:
  - Only the new rows appended in the result are written
  - Applicable only when existing rows in the result are not expected to change
- *Update Mode*:
  - Only outputs the rows that have changed since the last trigger
  - If the query doesn't contain aggregations, it will be equivalent to Append mode.

**Note** that each **mode** is applicable on certain types of queries.

# Example



Model of the Quick Example



# Notes

- Spark Structured Streaming does not materialize the entire table
  - reads the latest available data from the stream
  - processes it incrementally to update the result
  - discards the source data
- It only keeps around the minimal intermediate *state* data
  - e.g. intermediate counts

# API using Datasets and DataFrames

Since Spark 2.0, DataFrames and Datasets can represent static, bounded data, as well as streaming, unbounded data.

# Creating streaming DataFrames and streaming Datasets

Streaming DataFrames can be created through the `DataStreamReader` interface

([Scala/Java/Python docs](#))

returned by `SparkSession.readStream()`.

- Input Sources:
  - **File source** - Supported file formats: text, CSV, JSON, ORC, Parquet.
  - **Kafka source** - Reads data from Kafka (see the [Kafka Integration Guide](#))
  - **Socket source (for testing)** - Reads UTF8 text data from a socket connection. It does not provide end-to-end fault-tolerance guarantees.

# E.g., python

```
{% highlight python %}  
spark = SparkSession. ...
```

```
# Read text from socket  
socketDF = spark \  
    .readStream \  
    .format("socket") \  
    .option("host", "localhost") \  
    .option("port", 9999) \  
    .load()
```

```
socketDF.isStreaming()    # Returns True for DataFrames that have streaming sources
```

```
socketDF.printSchema()
```

```
# Read all the csv files written atomically in a directory  
userSchema = StructType().add("name", "string").add("age", "integer")  
csvDF = spark \  
    .readStream \  
    .option("sep", ";") \  
    .schema(userSchema) \  
    .csv("/path/to/directory") # Equivalent to format("csv").load("/path/to/directory")
```

# Basic Operations - Selection, Projection, Aggregation

— NOTE: **few operations** are not supported

```
df = ... # streaming DataFrame with IOT device data with schema { device: string, deviceType: string, signal: double, time: DateType }

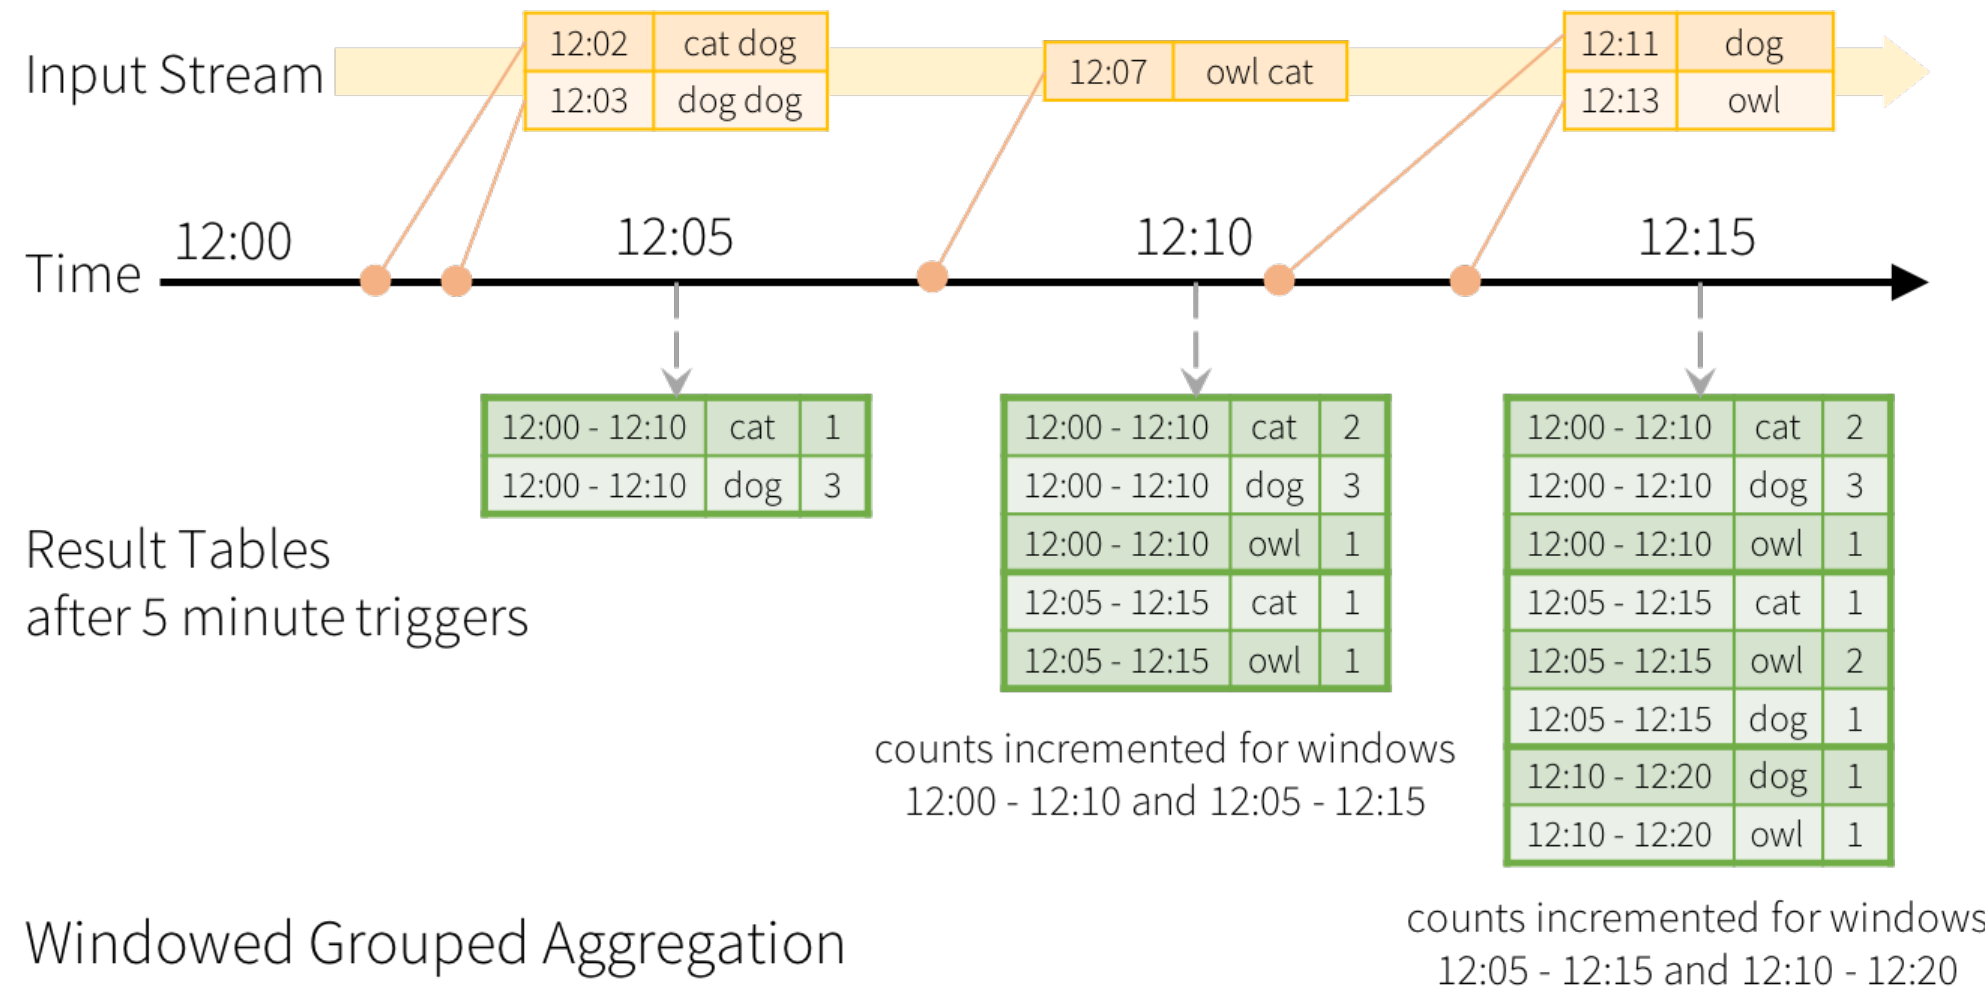
# Select the devices which have signal more than 10
df.select("device").where("signal > 10")

# Running count of the number of updates for each device type
df.groupBy("deviceType").count()
```

Alternatively, register a streaming DataFrame/Dataset as a temporary view and apply SQL on it

```
df.createOrReplaceTempView("updates")
spark.sql("select count(*) from updates") # returns another streaming DF
```

# Window Operations on Event Time



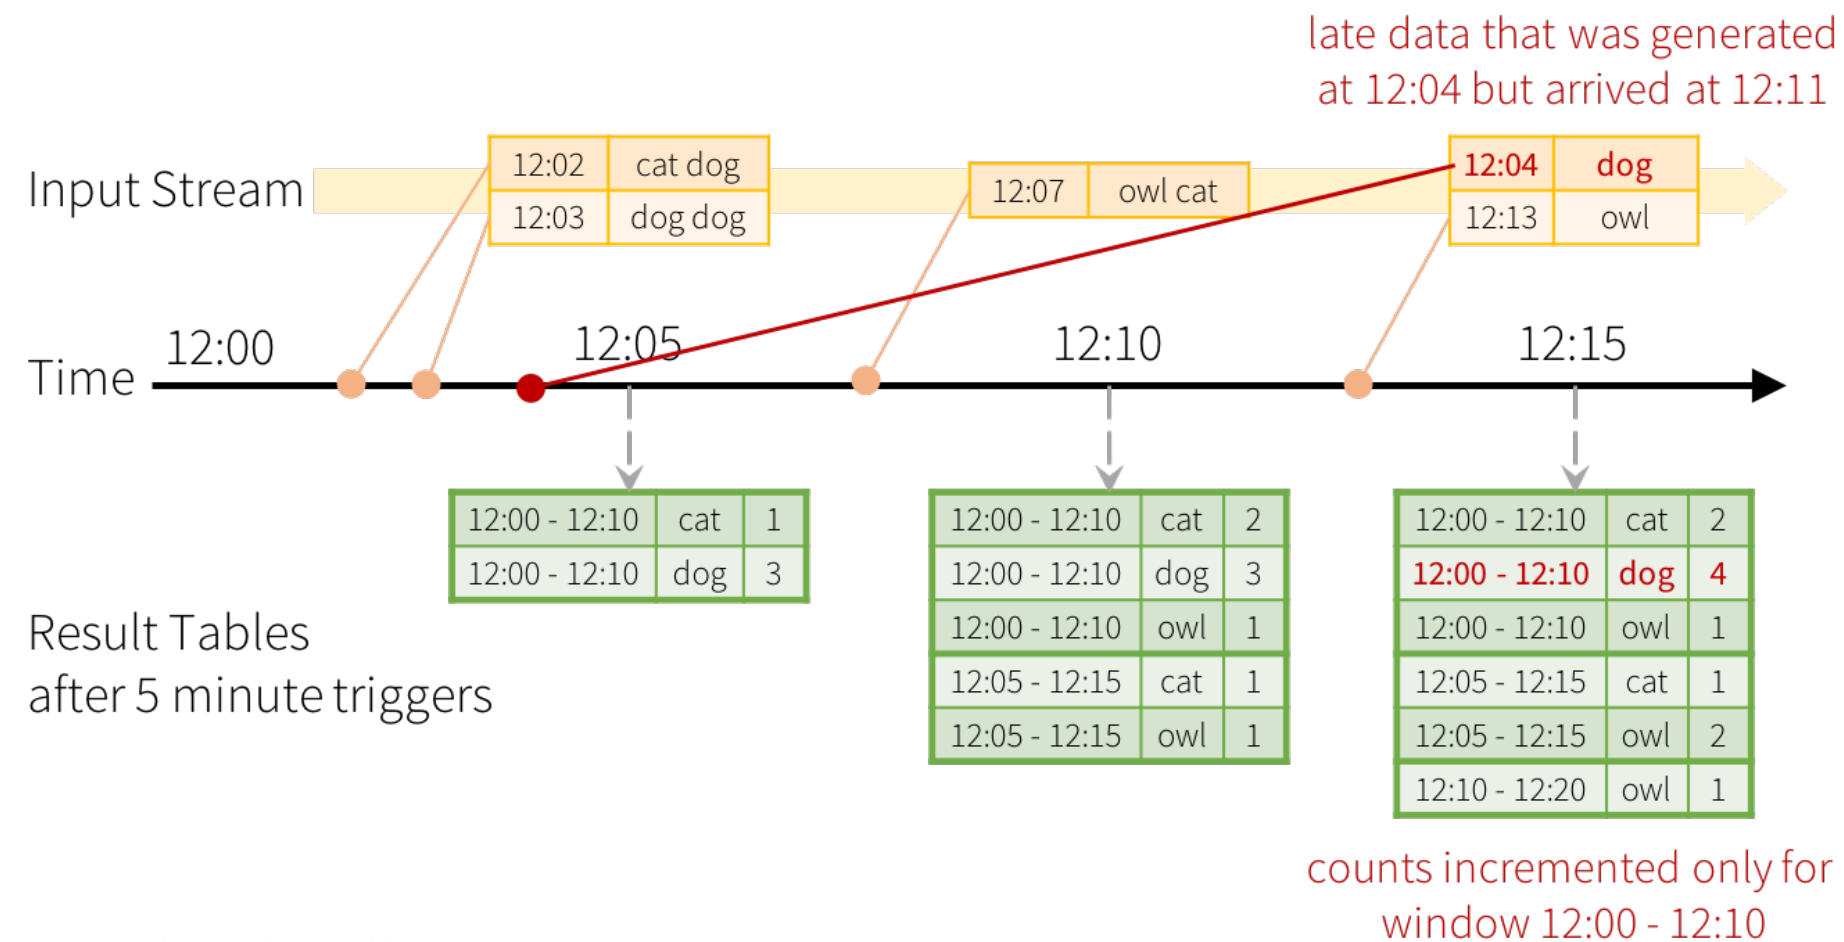
# Window Operations on Event Time (cont.)

Spark Structured Streaming treats this type of windows as grouping clauses.

```
words = ... # streaming DataFrame of schema { timestamp: Timestamp, word: String }

# Group the data by window and word and compute the count of each group
windowedCounts = words.groupBy(
    window(words.timestamp, "10 minutes", "5 minutes"),
    words.word
).count()
```

# Handling Late Data and Watermarking






# Handling Late Data and Watermarking (cont.)


```
words = ... # streaming DataFrame of schema { timestamp: Timestamp, word: String }

# Group the data by window and word and compute the count of each group
windowedCounts = words \
    .withWatermark("timestamp", "10 minutes") \
    .groupBy(
        window(words.timestamp, "10 minutes", "5 minutes"),
        words.word) \
    .count()
```

# Watermarking in Update Mode



# Watermarking in Append Mode



# Semantic Guarantees of Aggregation with Watermarking

- A watermark delay of "X hours" guarantees that the engine will never drop any data that is less than x hours delayed.
- **However**, data delayed by more than 2 hours is not guaranteed to be dropped
  - it may or may not get aggregated
  - the more delayed, the less likely

# Join Operations

The result of the streaming join is **generated incrementally**

# Stream-static Joins

```
staticDf = spark.read. ...  
streamingDf = spark.readStream. ...  
streamingDf.join(staticDf, "type") # inner equi-join with a static DF  
streamingDf.join(staticDf, "type", "left_outer") # left outer join with a static DF
```

## Notes

- INNER and LEFT-OUTER stream-static joins are not stateful
- OUTER and RIGHT-OUTER stream-static joins are not supported

for more information.

# Stream-stream Joins

- **Problem:** Any row received from one input stream can match with any future, yet-to-be-received row from the other input stream.
- **Solution:**
  - **Windows**, but not the kind you can express as a group by clause
    - buffer past input as streaming state
    - handle late, out-of-order data using watermarks

# Stream-stream Joins (cont.)

```
from pyspark.sql.functions import expr

impressions = spark.readStream. ...
clicks = spark.readStream. ...

# Apply watermarks on event-time columns
impressionsWithWatermark = impressions.withWatermark("impressionTime", "2 hours")
clicksWithWatermark = clicks.withWatermark("clickTime", "3 hours")

# Join with event-time constraints
impressionsWithWatermark.join(
    clicksWithWatermark,
    expr("""
        clickAdId = impressionAdId AND
        clickTime >= impressionTime AND
        clickTime <= impressionTime + interval 1 hour
        """))
)
```



# Stream-stream Joins

## Join Type

Inner	Supported, optionally specify watermark on both sides + time constraints for state cleanup
Left Outer	Conditionally supported, must specify watermark on right + time constraints for correct results, optionally specify watermark on left for all state cleanup
Right Outer	Conditionally supported, must specify watermark on left + time constraints for correct results, optionally specify watermark on right for all state cleanup
Full Outer	Not supported

## Policy for handling multiple watermarks

A streaming query can have multiple input streams that are unioned or joined together.

Each of the input streams can have a different threshold of late data that needs to

be tolerated for stateful operations. You specify these thresholds using `withWatermarks("eventTime", delay)` on each of the input streams. For example, consider

a query with stream-stream joins between `inputStream1` and `inputStream2`.

</div>

While executing the query, Structured Streaming individually tracks the maximum

# Starting Streaming Queries

- Once you have defined the final result DataFrame/Dataset, all that is left is for you to start the streaming computation.
- To do that use the DataStreamWriter  
([Scala/Java/Python docs](#))  
returned through `Dataset.writeStream()`

# Starting Streaming Queries (cont.)

- specify:
  - *output sink*: Data format, location, etc.
  - *Output mode*: what gets written
  - *Query name*: Optionally, a unique name of the query
  - *Trigger interval*: Optionally, the trigger interval
  - *Checkpoint location*: for the end-to-end fault-tolerance
- Output Sinks:
  - **File sink** - a directory
  - **Kafka sink** - one or more topics in Kafka
  - **Foreach sink** - Runs arbitrary computation
  - **Console sink (for debugging)**

# Managing Streaming Queries

```
query = df.writeStream.format("console").start()    # get the query object

query.id()          # get the unique identifier of the running query that persists across restarts from checkpoint data

query.runId()       # get the unique id of this run of the query, which will be generated at every start/restart

query.name()        # get the name of the auto-generated or user-specified name

query.explain()     # print detailed explanations of the query

query.stop()        # stop the query

query.awaitTermination() # block until query is terminated, with stop() or with error

query.exception()    # the exception if the query has been terminated with error

query.recentProgress() # an array of the most recent progress updates for this query

query.lastProgress() # the most recent progress update of this streaming query
```

# Further Reading

- See and run the [Scala/Java/Python/R](#) examples.
  - [Instructions](#) on how to run Spark examples
- Read about integrating with Kafka in the [Structured Streaming Kafka Integration Guide](#)
- Read more details about using DataFrames/Datasets in the [Spark SQL Programming Guide](#)
- Third-party Blog Posts
  - [Real-time Streaming ETL with Structured Streaming in Apache Spark 2.1 \(Databricks Blog\)](#)
  - [Real-Time End-to-End Integration with Apache Kafka in Apache Spark's](#)

# Talks

- Spark Summit Europe 2017
  - Easy, Scalable, Fault-tolerant Stream Processing with Structured Streaming in Apache Spark - [Part 1 slides/video](#), [Part 2 slides/video](#)
  - Deep Dive into Stateful Stream Processing in Structured Streaming - [slides/video](#)
- Spark Summit 2016
  - A Deep Dive into Structured Streaming - [slides/video](#)

# Spark Structured Streaming

---

**Emanuele Della Valle**

---

prof @ Politecnico di Milano  
Co-founder @ Quantia Consulting

[emanuele.dellavalle@polimi.it](mailto:emanuele.dellavalle@polimi.it)

<http://emanueledellavalle.org>