

ksqlDB

Emanuele Della Valle

prof @ Politecnico di Milano

Why use ksqlDB?

Why use ksqlDB?

Too many cooks in the kitchen.

Why use ksqlDB?

Divide and conquer

Why use ksqlDB?

Cut from the same mold

Why use ksqlDB?

Reunited

Key operations

Key operations

A running example about riders

Key operations

Capture events

```
CREATE SOURCE CONNECTOR riders WITH (  
  'connector.class' = 'JdbcSourceConnector',  
  'connection.url'  = 'jdbc:postgresql://...',  
  'topic.prefix'    = 'rider',  
  'table.whitelist' = 'riderLocations, profiles',  
  'key'             = 'profile_id',  
  ...);
```

Key operations

Perform continuous transformations

```
CREATE STREAM locations AS
  SELECT rideId, latitude, longitude,
         GEO_DISTANCE(latitude, longitude,
                       dstLatitude, dstLongitude
                       ) AS kmToDst

  FROM ridersLocation
  EMIT CHANGES;
```

// Enrichment

Key operations

Create materialized views

```
CREATE TABLE activePromotions AS  
  SELECT rideId,  
         qualifyPromotion(kmToDst) AS promotion  
FROM locations  
GROUP BY rideId  
EMIT CHANGES;
```

User Defined Agg

Key operations

Serve lookups against materialized views

```
SELECT rideId, promotion  
FROM activePromotions  
WHERE rideId = '6fd0fcdb';
```

Concepts

Concepts

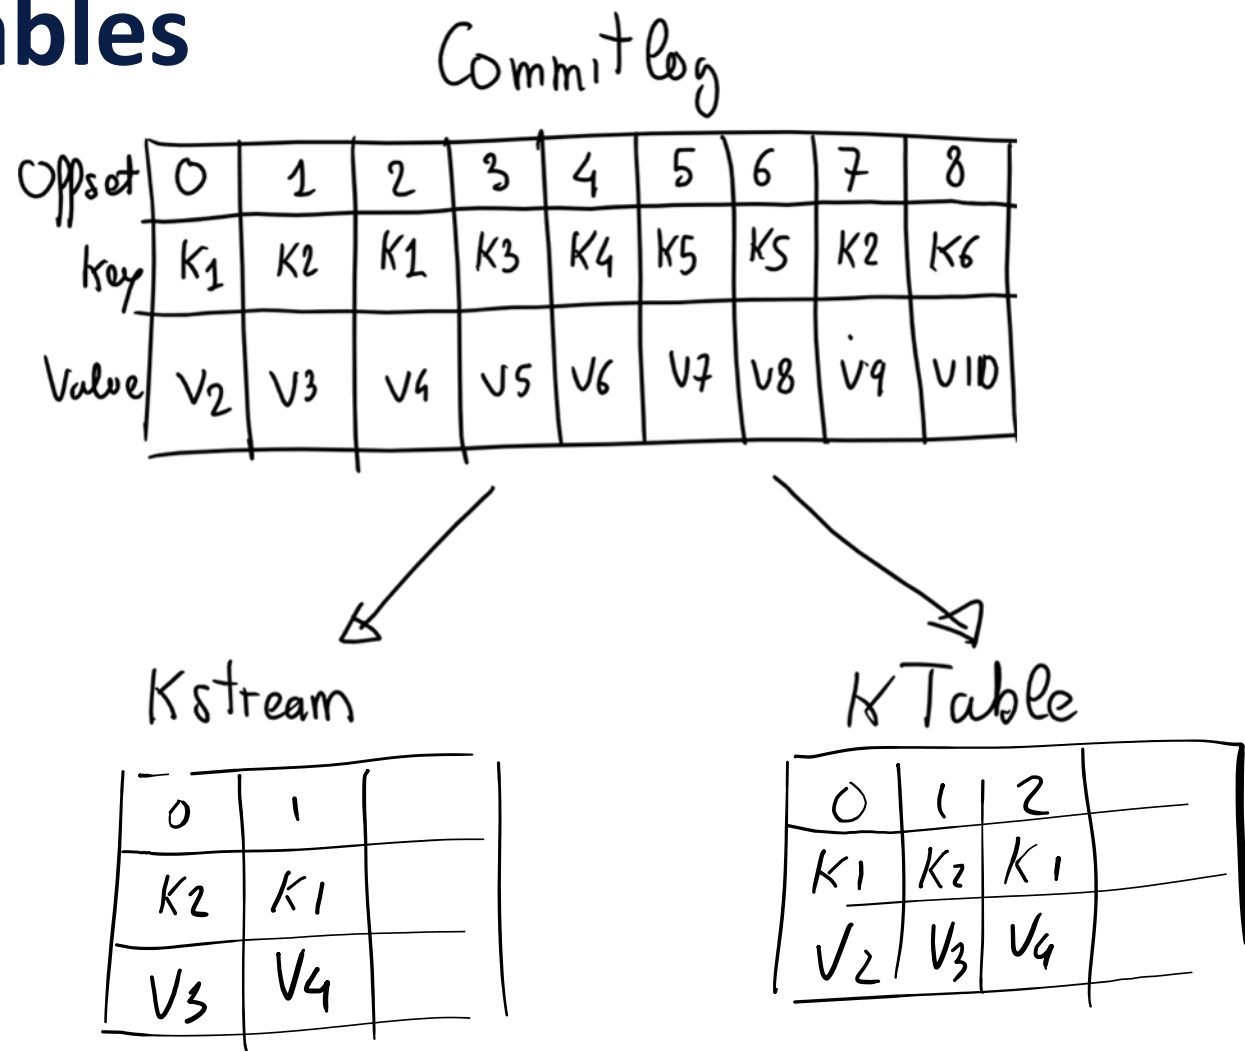
Basics

- **Event**: the fundamental unit of data in stream processing
- **Stream**: an *immutable, append-only** collection of events that represents a series of historical facts
- **Table** (a.k.a. **materialized view**): a *mutable* collection of events. They let you represent the latest version of each value per key.
- **Queries**: how you *transform, filter, aggregate, and join* collections together to derive new collections or materialized views that are incrementally updated in real-time as new events arrive.

Concepts

Topics vs. Streams vs. Tables

- topics can have two origins in Kafka
 - *independent observations*
 - *change logs* captured from a DB
- *independent observations* can be captured in a **stream**
- *change logs* can be captured in a **stream** but can also in a time-varying **table** (a.k.a. **materialized views**)



Concepts

Push queries

- let you **subscribe** to a query's result as it changes in real-time
- When new events arrive, push queries emit refinements, which allow **reacting** to new information
- Fit for **asynchronous** application flows

```
SELECT riderId, latitude, longitude
FROM Locations
WHERE rider = '6fd0fcdb'
EMIT CHANGES;
```


Concepts

Pull queries

- **Fetch** the current state of a materialized view
- Incrementally updated as new events arrive
- fit for **request/response flows**

```
SELECT riderId, latitude, longitude  
FROM currentCarLocations  
WHERE ROWKEY = '6fd0fcdb';
```

How it works

How it works

ksqlDB Architecture

How it works

ksqlDB vs Kafka Streams 1/3

How it works

ksqlDB vs Kafka Streams 2/3

- For example, the following KSQL query ...

→ create at registration time

```
CREATE STREAM fraudulent_payments AS
```

```
SELECT fraudProbability(data)
```

```
FROM payments
```

```
WHERE fraudProbability(data) > 0.8
```

```
EMIT CHANGES;
```

UDF

Filter

src

How it works

ksqlDB vs Kafka Streams 3/3

... is equivalent to the following Scala code

```
object FraudFilteringApplication extends App {  
  val builder: StreamsBuilder = new StreamsBuilder()  
  val fraudulentPayments: KStream[String, Payment] = builder  
    .stream[String, Payment]("payments")  
    .filter((_, payment) => payment.fraudProbability > 0.8)  
    fraudulentPayments.to("fraudulent-payments-topic") UPF  
  val config = new java.util.Properties  
  config.put(StreamsConfig.APPLICATION_ID_CONFIG, "fraud-filtering-app")  
  config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "kafka-broker1:9092")  
  val streams: KafkaStreams = new KafkaStreams(builder.build(), config)  
  streams.start()  
}
```

How it works

Differences Between ksqlDB and Kafka Streams

What	ksqlDB	Kafka Streams
You write	<i>~ SQL</i>	<i>scala / java</i>
Console	✓	✗
REST API	✓	✗
Runtime	✓ <i>ksqlDB server</i>	<i>JVM</i>

How it works

Deployment Modes

How it works

Fault tolerance and scale-out

Notes

- For each query there is a src topic (with its partitions, say 3)
- Each ksqlDB server creates a consumer referring to the consumer group of the query
- As far as there are partitions to assign, Kafka assigns them to the consumers created by the ksqlDB Serves
 - Up to 3 in the example
- Any other consumer created by the ksqlDB Serves will remain in idle state waiting for a fault or for a scale out operation

How it works

Query Lifecycle

Walking through basic tutorials

Walking through basic tutorials

Materialized Views

The hard way

The ksqlDB way

Walking through basic tutorials

Streaming ETL

The hard way

The ksqlDB way

Learn more

- <https://ksqldb.io/>
- <https://ksqldb.io/examples.html>
- <https://github.com/confluentinc/ksql>
- <https://twitter.com/ksqldb>

Thank you for your attention!

Questions?

ksqlDB

Emanuele Della Valle

prof @ Politecnico di Milano