

**Streaming Data Analytics**  
**Preview of the**  
**part of the exam about**  
**Streaming Data Engineering**

Emanuele Della Valle  
Politecnico di Milano



**POLITECNICO**  
MILANO 1863

# Exam content

- **Questions** on all lectures about Streaming Data Engineering **to test**

- **The breadth** of your knowledge
- **The depth** of your knowledge

(8 points)

- **Exercises**

- Given a sequence of events and an **EPL pattern** using the followed-by operator (->) and pattern **guards**, tell **which** events the pattern matches, which it does not, and **why**
- Given a realistic **Streaming Data Engineering problem**, **define** the **schema** of the streams, **provide** a minimal example of the **events** in those streams, **propose continuous queries** and **show** the **results** you expect from them. Formulate the solution **in EPL** and **in a language you choose between** KSQL and Spark Structured Streaming. **Explain** your code in detail.

(7 points)

# Questions to test the **breadth** of your knowledge

- Which are the three approaches to tame velocity? List them all before comparing and contrasting two of them
- Which are the three time-models we introduced? List them all before comparing and contrasting two of them
- What are the types of windows? List them all and describe the behavior of at least two of them
- Illustrate the programming model of spark structured streaming at the logical and physical level
- Describe Kafka at a conceptual, system, and physical level

# Questions to test the **breadth** of your knowledge

- How does ksqlDB work? Illustrate the architecture of a ksqlDB server and explain how it coordinates with other ksqlDB servers
- How is a join performed between a stream and a table? Given an example that illustrates the different results you obtain using an inner, left outer, right outer, or outer join
- Under which constraint is it possible to join two streams? Given an example that illustrates the different results you obtain using an inner or a left outer join
- How does Spark Structured Stream treat late arrivals?

# Questions to test the **depth** of your knowledge

- What's the difference between the stream-only and the absolute time models? Explain it, proposing examples of queries that can be answered in both models and queries that can be answered only with the absolute time model.
- What's the difference between the absolute and the interval-based time models? Explain it, proposing examples of queries that can be answered in both models and queries that can be answered only with the interval-based time model.
- What are streams and events? How do they relate to each other? Give an example.
- What is a tubling logical window? Give an example at the conceptual level and show that you know both the EPL syntax and that of another language of your choice.

# Questions to test the **depth** of your knowledge

- What is a hopping logical window? Give an example at the conceptual level and show that you know both the EPL syntax and that of another language of your choice.
- What is the role of the output clause in EPL? Give an example that supports your explanation.
- What's the role of watermarking in Spark Structured Streaming? How does it interleave with the output mode? Support your claims with an example.

# Questions to test the **depth** of your knowledge

- What's the role of a topic in Kafka? How does it relate to partitions and brokers? Explain it using an example.
- What's the role of a broker in Kafka? How does it relate to topics and partitions? Explain it using an example.
- What's the role of a consumer group in Kafka? How does it relate to topics and partitions? Explain it using an example.
- What is a session window? Give an example at the conceptual level and show that you know the ksqlDB syntax

# Exercises on EPL patterns

- Suppose you receive the following stream of events:

`A1@0, C1@1, B1@2, B2@3, A2@4, B3@5, A3@6, B4@10.`

- Note that `A3@6` denotes an event of type `A` identified by the number `3` that is received at time `6`.
- Given the patten:

`every A -> (B and not C where timer:within(3 sec))`

- Translate the pattern into an English sentence
- Which are the events that trigger the matching? Why?
- Which are the events that may trigger the matching but are excluded by the semantics of the `every` and the `where timer:within` clauses? Why?



# Exercises on a Streaming Data Engineering problem

- In the git repo of the course you find a complete example about a Robotic Arm solved in
  - EPL
    - [https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl\\_robotic-arm](https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl_robotic-arm)
  - ksqlDB
    - [https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/ksql\\_robotic-arm](https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/ksql_robotic-arm)
  - Spark Structured Streaming
    - [https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/ssss\\_robotic-arm](https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/ssss_robotic-arm)
- There are other two example completely solved in EPL
  - [https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl\\_tomatopick](https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl_tomatopick)
  - [https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl\\_bocce](https://github.com/emanueledellavalle/streaming-data-analytics/tree/main/codes/epl_bocce)

Streaming Data Analytics

**Preview of the  
part of the exam about  
Streaming Data Engineering**

Emanuele Della Valle  
Politecnico di Milano



**POLITECNICO**  
MILANO 1863